

Approximating Boosted Decision Trees with Differential Privacy


Thorsten Peinemann

My personal website: tpein160.github.io




UNIVERSITÄT ZU LÜBECK
INSTITUTE FOR IT SECURITY
PRIVACY & SECURITY GROUP

Temperature prediction



Sun in the last hour, yes/no?



Current temperature in celsius



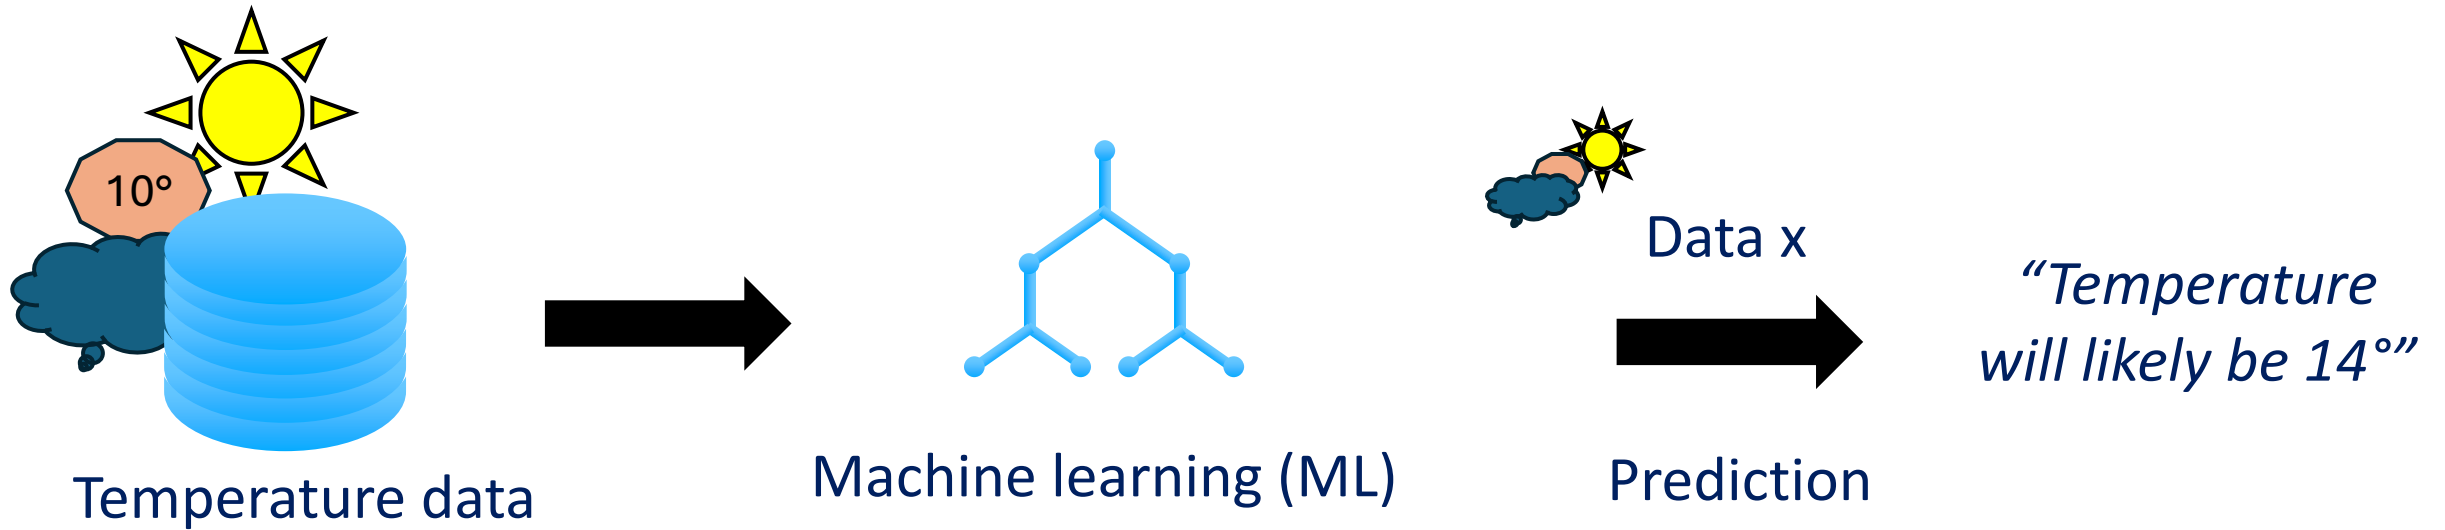
Rain in the last hour, yes/no?



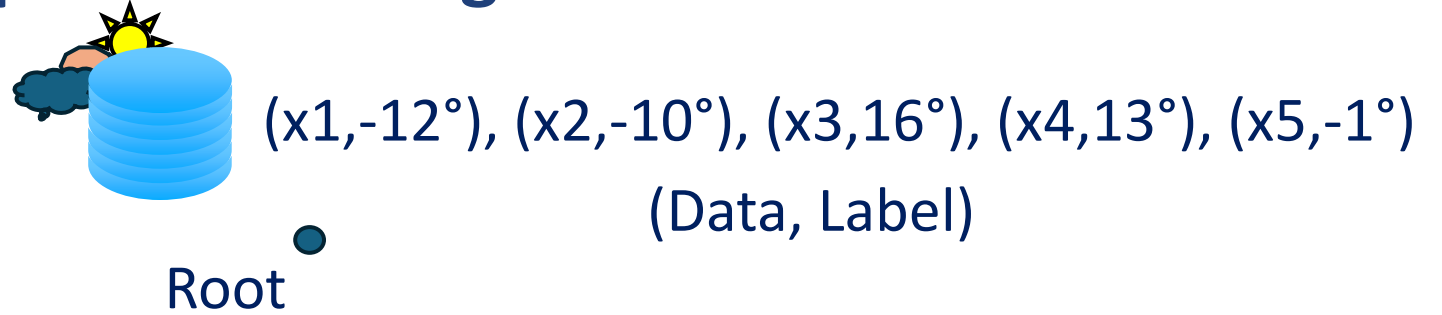
Predict temperature
in 3 hours




Temperature prediction using boosted decision tree (BDT) model



Non-private training of BDT model




Non-private training of BDT model




Each node splits the data in two subsets so that each subset groups together alike labels (e.g. gini-coefficient)

Non-private training of BDT model




Each node splits the data in two subsets so that each subset groups together alike labels (e.g. gini-coefficient)

Non-private training of BDT model

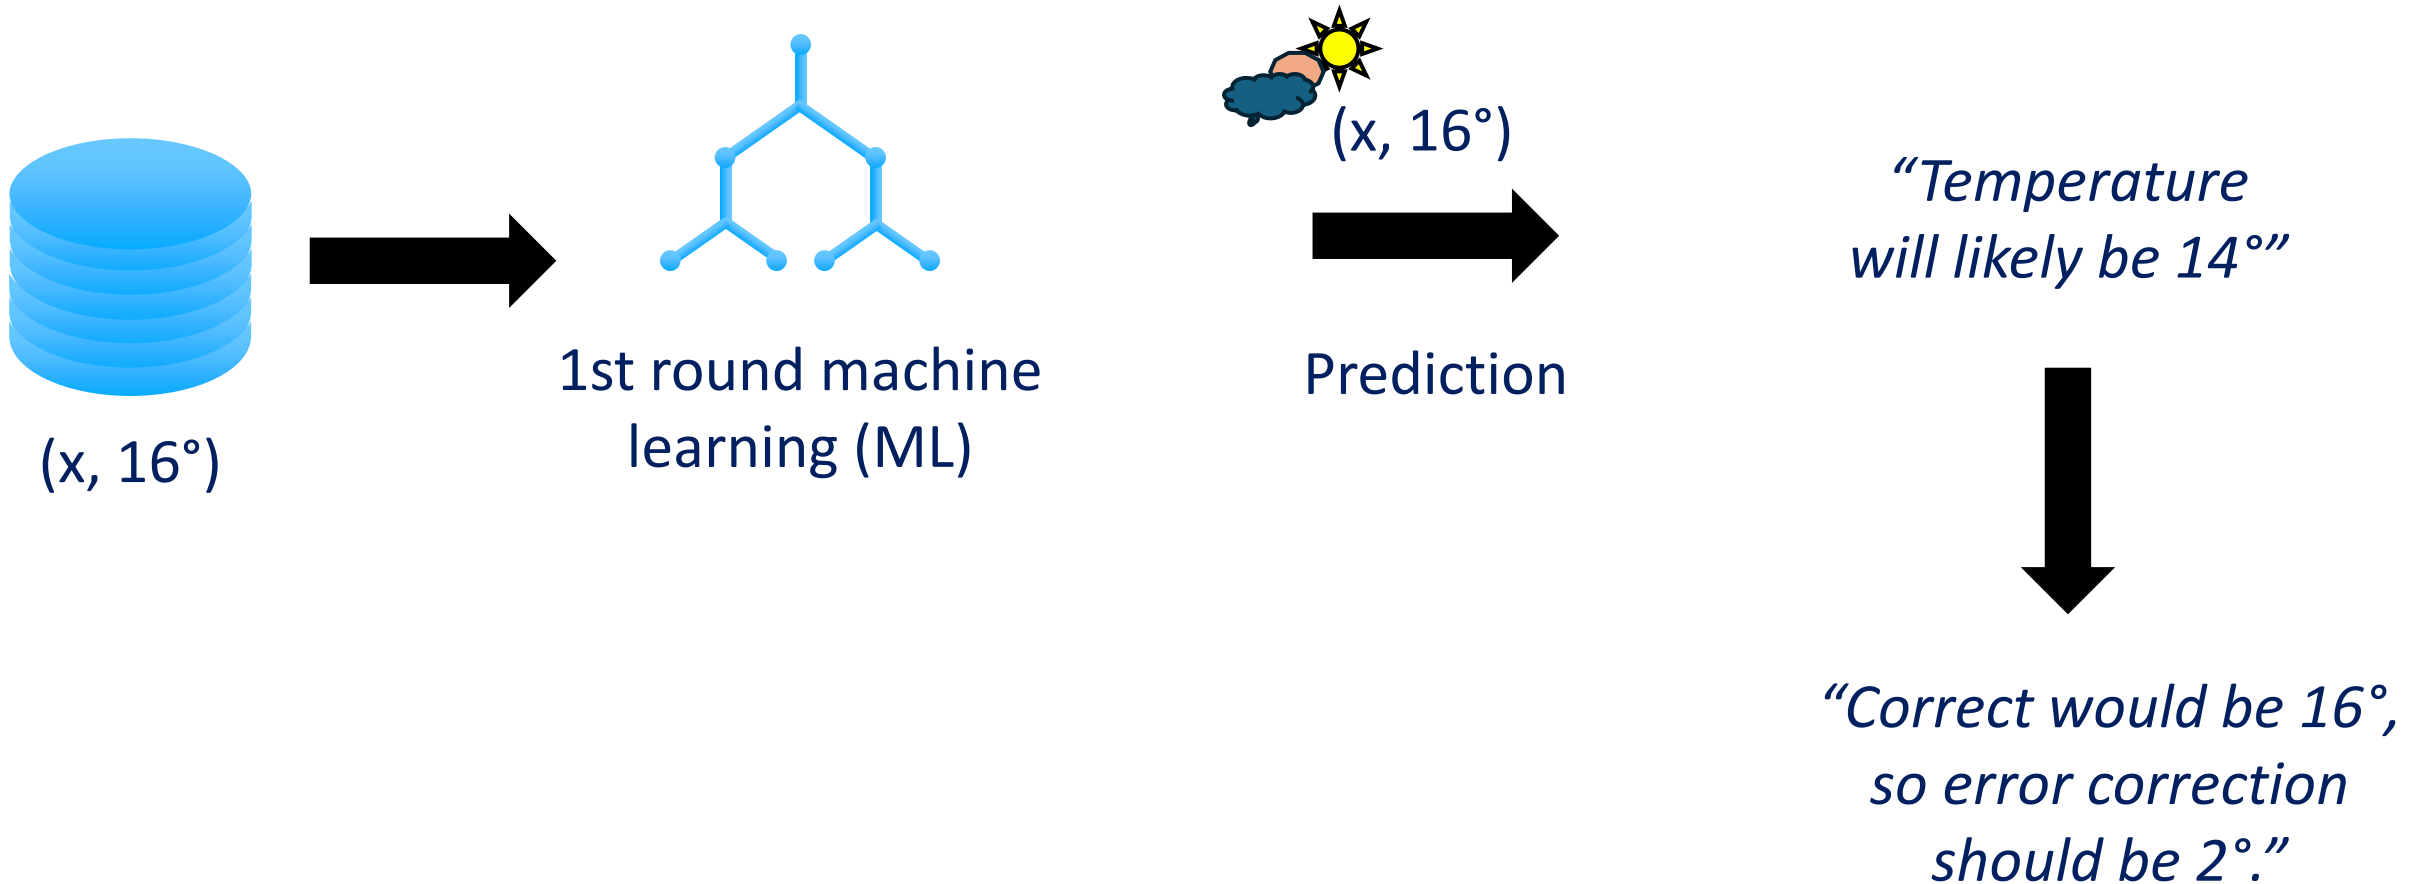


A leaf stores the average label of data points in that leaf

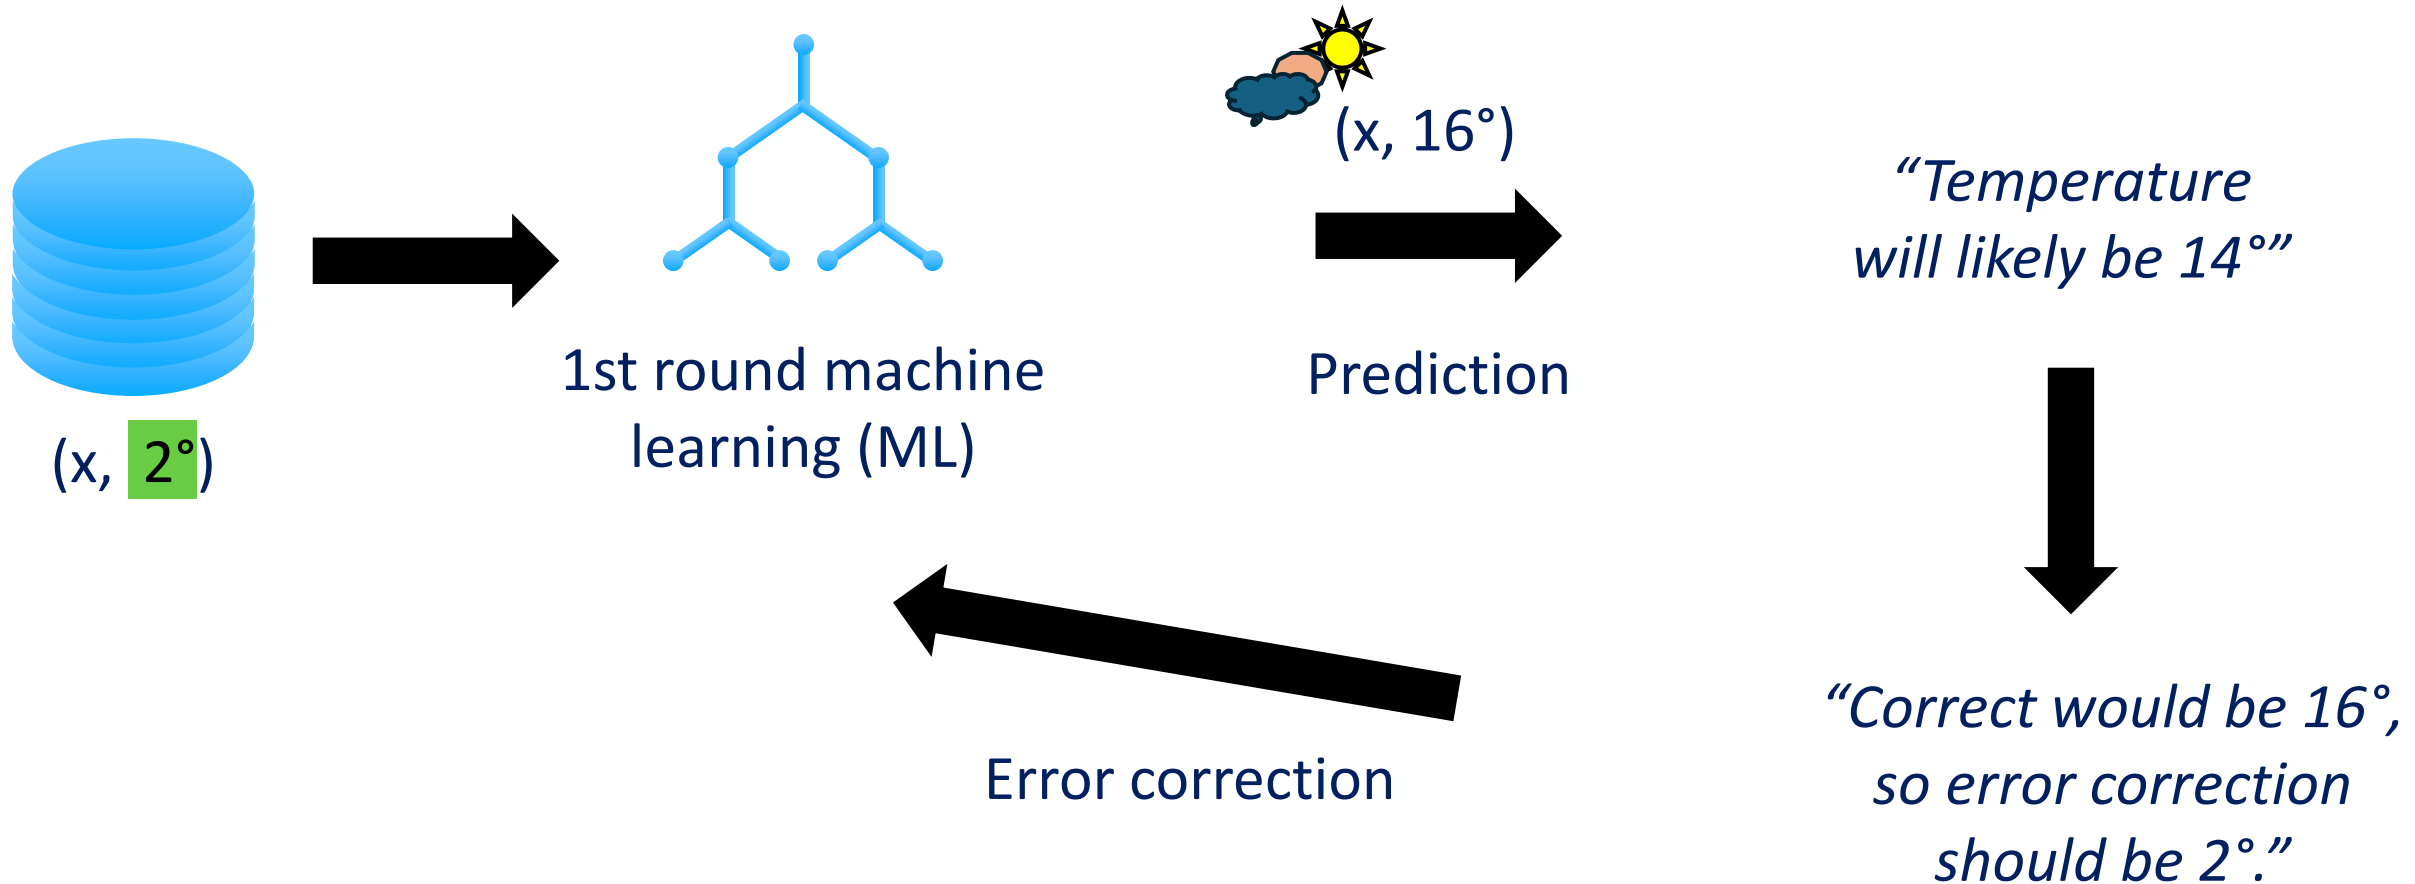
Prediction of BDT model



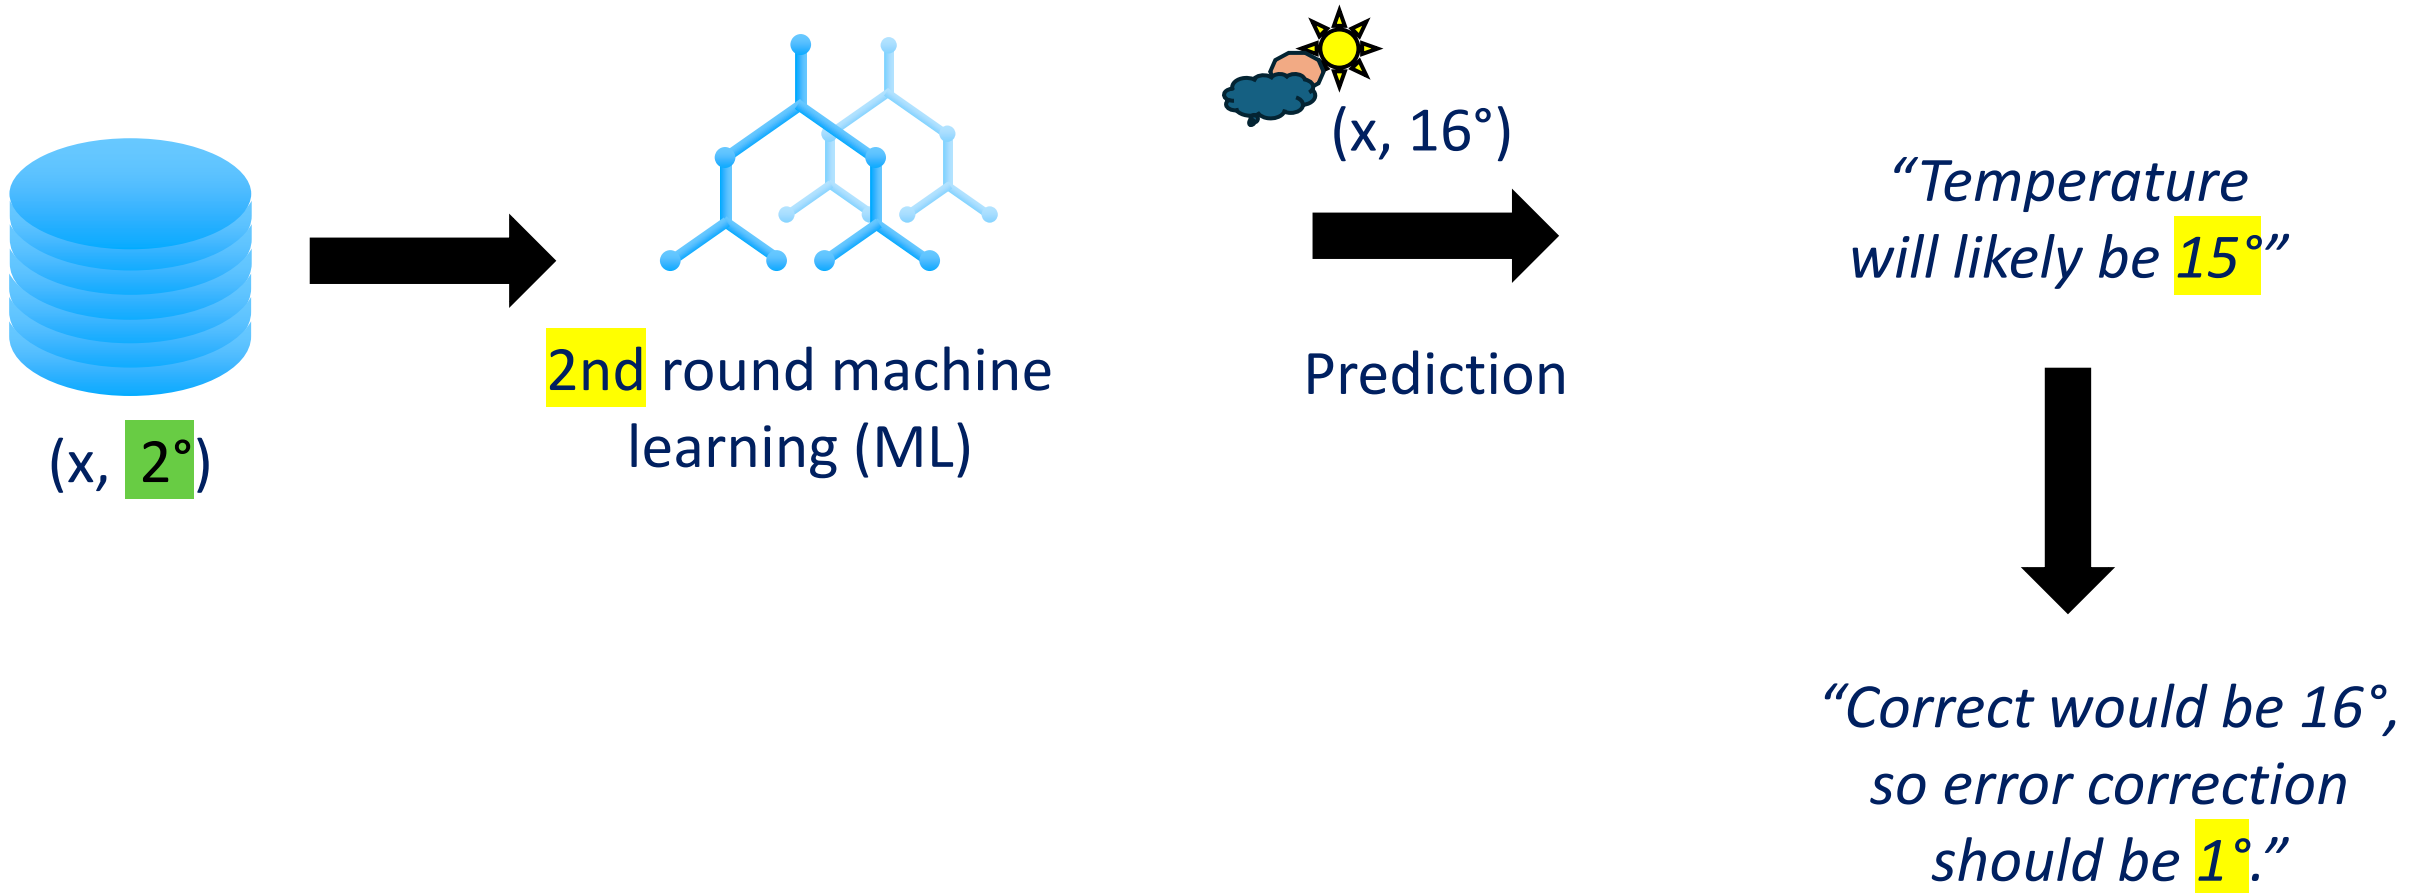
Error correction for iterative BDT training



Error correction for iterative BDT training



Error correction for iterative BDT training




Differentially private boosted decision trees

Whence cometh the noise?



Data-dependent part of BDT model



Algorithms for splitting and leaves

Splitting

From top to bottom:

- Split each node until maximum depth
- Split such that equal error-corrected labels are grouped

Leaves

For each leaf:

- Find all data points in this leaf
- Sum their error-corrected labels $\rightarrow S$
- Count number of data points $\rightarrow C$
- Store S/C



DP-approximated splitting algorithm

DP-Splitting

From top to bottom:

Split each node until max depth is reached

Split randomly

Leaves

For each leaf:

Find all data points in this leaf

Sum their error-corrected labels $\rightarrow S$

Count number of data points $\rightarrow C$

Store S/C



DP-approximated splitting and leaves algorithms

DP-Splitting

From top to bottom:

Split each node until max depth is reached

Split randomly

DP-Leaves

For each leaf:

Find all data points in this leaf

Clip their error-corrected labels to length L

Sum clipped error-corrected labels $\rightarrow S_c$

Add Gaussian noise: $S_c \rightarrow S'_c$

Count number of data points $\rightarrow C$

Add Gaussian noise: $C \rightarrow C'$

Store S'_c/C'



DP-Proof for DP-approximated splitting and leaves algorithms

DP-Splitting

Output of randomized function has no leakage



DP-Proof for DP-approximated splitting and leaves algorithms

DP-Splitting

Output of randomized function has no leakage

DP-Leaves

(ϵ, δ) -Differential Privacy (DP):

$$\Pr[M(D) \in S] \leq e^\epsilon \Pr[M(D \cup \{x\}) \in S] + \delta$$

DP-Leaves:

(1) Leakage for x occurs only in x 's leaf \mathbf{P}_x

$$(2) M(D) = \left(\sum_{(v,l) \in D: (v,l) \in \mathbf{P}_x} \text{clip}(l, (-L, L)) \right) + N(0, \sigma^2)$$

Gaussian Mechanism:

M satisfies (ϵ, δ) -DP for any $\delta > 0$, ϵ in $(0, 1)$ when

$$\sigma > \sqrt{2 \ln(1.25/\delta)} * L/\epsilon$$



Approximating Boosted Decision Trees with Differential Privacy

Thorsten Peinemann

My personal website: tpein160.github.io



UNIVERSITÄT ZU LÜBECK
INSTITUTE FOR IT SECURITY
PRIVACY & SECURITY GROUP